HANUMANTHA RAO VULLI

Denton, TX | vullihanumantharao@gmail.com
Portfolio: https://sunnyvulli.github.io/ | GitHub: https://github.com/Sunnyvulli

Dear Hiring Manager,

I am seeking a data engineering role where I can contribute to reliable pipelines, usable analytics, and maintainable backend services. I hold a Master of Science in Data Engineering from the University of North Texas. My professional background combines AWS data workflows, SQL, and Java development with Spring Boot.

At PRC Cloud Tech, I contribute to ingestion and ETL workflows, write and optimize SQL, support data validation, and help troubleshoot scheduled processing failures. Earlier roles involved PostgreSQL and Oracle query tuning, Spring Boot services integrated with AWS, and Jenkins and Docker delivery workflows. These responsibilities have reinforced the value of clear data checks and documentation that other engineers can use.

My portfolio provides inspectable examples. Lakehouse Guardian processes 17,379 UCI Bike Sharing observations with Spark, local Iceberg snapshots, dbt checks, and a LangGraph approval gate. Its real Airflow integration and Terraform S3/Glue configuration pass CI; AWS infrastructure has not been deployed. Across 21 controlled fault batches, row detection precision was 1.00 and recall 0.80, with missed plausible count inflation documented. Catalog Sentinel adds read-only MCP tools with strict schemas, host-assigned permissions, and audit records: its actual stdio evaluation enforced all 13 expected denials in 16 requests, with database bytes unchanged.

Evidence Ledger RAG adds incremental indexing and cited extracts with 18 tests: hybrid expected-file hit@3 was 11/14 versus BM25 10/14 on 14 authored queries, with semantic guard failures documented. Mobility Quality Lab includes 13 passing tests and a reproducible historical backtest. My Applied ML Research Suite covers citation-support triage, phishing robustness, and simulated-transient ranking, with 58 integrity tests and documented negative findings. These repositories distinguish local engineering evidence, controlled experiments, and unpublished writing from production validation.

I would welcome the opportunity to contribute my data and backend engineering experience to your team. I value collaborative debugging, clear technical communication, and practical improvements that make data dependable. My resume and portfolio provide the supporting detail, and I would be glad to discuss the requirements of a specific role.

Thank you for your consideration.

Sincerely,
Hanumantha Rao Vulli